

50.007 Machine Learning

2026 Spring

2. Perceptron

Na Zhao
Assistant Professor, ISTD

Recap – Steps to solving a supervised learning problem

1. Decide what the input-output pairs are.
2. Decide how to encode inputs and outputs.
 - This defines the input space \mathcal{X} , and the output space \mathcal{Y} .
3. Choose a *class of hypotheses \mathcal{H}* , a.k.a. the *model* or *set of classifiers*.
4. Choose an *error function (cost function)* to define the best hypothesis h .
5. Choose a *learning algorithm/criterion* for searching efficiently through the space of hypotheses - *optimization*.


Recap – Linear Classifier

Linear classifier is a **highly constrained** set of classifiers (through origin):

$$h(x; \theta) = \text{sign}(\theta_1 x_1 + \dots + \theta_d x_d) = \text{sign}(\theta \cdot x) = \begin{cases} +1, & \theta \cdot x \geq 0 \\ -1, & \theta \cdot x < 0 \end{cases}$$

- $\theta = [\theta_1, \dots, \theta_d]^T$ is a column vector of real valued **parameters** or **weights**
- Different settings of the **weights** give rise to **different classifiers**.

theta here is normal vector



- A **decision boundary** is a **line** (or a **surface**, in higher dimensions) that **separates** the two classes, resulting in two half-spaces:
$$\mathcal{X}^+(\theta) = \{x \in \mathcal{R}^d : h(x; \theta) = +1\}$$
$$\mathcal{X}^-(\theta) = \{x \in \mathcal{R}^d : h(x; \theta) = -1\}$$
- Note that these half spaces as well as the decision boundary, depend on **how we set θ** .

Learning Objectives

You should be able to know:

1. What is a linearly separable set of examples?
2. How does the Perceptron algorithm work?
3. What is the guarantee of the Perceptron algorithm when the dataset is linearly separable?

Linear Classifier

$$h(x; \theta) = \text{sign}(\theta_1 x_1 + \dots + \theta_d x_d) = \text{sign}(\theta \cdot x) = \begin{cases} +1, & \theta \cdot x \geq 0 \\ -1, & \theta \cdot x < 0 \end{cases}$$

h - prediction
y - true value

- How should we pick θ ? - choose theta chooses final model

Error minimization!

- Intuitively, θ should make the predictions of h close to the true values y on the data we have (training data set)
- Hence, we will define an **error function** or **cost function** to measure how much our prediction differs from the “true” answer:

$0.x(t) = 0, 0 \text{ lies on the decision boundary}$

$$\mathcal{E}_n(\theta) = \frac{1}{n} \sum_{t=1}^n [\![y^{(t)} \neq h(x^{(t)}; \theta)]\!] = \frac{1}{n} \sum_{t=1}^n [\![y^{(t)}(\theta \cdot x^{(t)}) \leq 0)]\!] \quad \leftarrow \text{zero-one Loss}$$

$[\![\cdot]\!]$ returns 1 if the logical expression in the argument is true, and zero otherwise.

n training sample and input x training data, h will be prediction data, y will be actual data


y can only have 2 values since binary classification,
true label is 1/-1

Error Minimization

- What would a reasonable algorithm be for finding the optimal θ that minimizes the training error?
 - Generally, this is **not an easy problem** to solve...
- But in this lecture, we consider **a special case** where there exists a **linear classifier** that achieves **zero training error**, known as the ***realizable*** case.
- “**Realizability**” depends on both the training examples as well as the set of classifiers we have adopted.
- Hence, we make two assumptions:
 1. The training examples are ***linearly separable through origin***.
 2. The hypothesis is a ***linear classifier through origin***.

Linearly Separable

Definition 1.1 Training examples $S_n = \{(x^{(t)}, y^{(t)}), t = 1, \dots, n\}$ are linearly separable through origin if there exists a parameter vector $\hat{\theta}$ such that $y^{(t)}(\hat{\theta} \cdot x^{(t)}) > 0$ for all $t = 1, \dots, n$.



linearly separable through origin



not linearly separable

The Perceptron Algorithm

```
Random initialize  $\theta$ 
while TRUE do
     $m = 0$    m - indicator to see how many run cases
    for  $(x^{(t)}, y^{(t)}) \in D$  do
        if  $y^{(t)}(\theta \cdot x^{(t)}) \leq 0$  then Perceptron update rule
             $\theta \leftarrow \theta + y^{(t)}x^{(t)}$ 
             $m \leftarrow m + 1$ 
        end if
    end for
    if  $m = 0$  then
        break
    end if
end while

        // Initialize  $\theta$ .  $\theta = \vec{0}$  misclassifies everything.  

        // Keep looping  

        // Count the number of misclassifications,  $m$   

        // Loop over each (data, label) pair in the dataset,  $D$   

        // If the pair  $(x^{(t)}, y^{(t)})$  is misclassified  

        // Update the weight vector  $\theta$   

        // Counter the number of misclassification
    

        // If the most recent  $\theta$  gave 0 misclassifications  

        // Break out of the while-loop
    

        // Otherwise, keep looping!
    
```

Adpted from Kilian Q. Weinberger's slide (Cornell University CS4780)

Perceptron Update Rule

How should we understand the perceptron update rule?



$$\text{if } y^{(t)} \neq h(x^{(t)}; \theta^{(k)}) \text{ then}$$
$$\theta^{(k+1)} = \theta^{(k)} + y^{(t)}x^{(t)}$$

- The weights θ are updated only if we make a mistake, and these updates do tend to correct mistakes.
- If the classifier predicted an example that was negative but it should have been positive...
 - Currently: $\theta \cdot x < 0$
 - Want: $\theta \cdot x > 0$
- Adjust the weights θ so that this function values becomes more positive.

Perceptron Update Rule – An Example

- The points $x^{(1)}$ and $x^{(2)}$ in the figure are chosen such that the algorithm makes a mistake on both of them during its first pass.
- $\theta^{(0)} = 0 \rightarrow \theta^{(1)} = \theta^{(0)} + x^{(1)} \rightarrow \theta^{(2)} = \theta^{(1)} + (-1)x^{(2)}$

Perceptron update rule:
if $y^{(t)} \neq h(x^{(t)}; \theta^{(k)})$ then
 $\theta^{(k+1)} = \theta^{(k)} + y^{(t)}x^{(t)}$



Perceptron Update Rule - Algebraical View

- Suppose we **make a mistake** on $x^{(t)}$
- Then the **updated weights** are given by $\theta^{(k+1)} = \theta^{(k)} + y^{(t)}x^{(t)}$
- If we consider **classifying the same example $x^{(t)}$** after the update, **using the new weights $\theta^{(k+1)}$**

$$\begin{aligned}y^{(t)}(\theta^{(k+1)} \cdot x^{(t)}) &= y^{(t)}(\theta^{(k)} + y^{(t)}x^{(t)}) \cdot x^{(t)} \\&= y^{(t)}(\theta^{(k)} \cdot x^{(t)}) + (y^{(t)})^2(x^{(t)} \cdot x^{(t)}) \\&= y^{(t)}(\theta^{(k)} \cdot x^{(t)}) + \boxed{\|x^{(t)}\|^2}\end{aligned}$$

- As a result of the update, the value of $y^{(t)}(\theta \cdot x^{(t)})$ increases (becomes more positive)
- If we consider the same example repeatedly, then we will necessarily change the weights such that the example will be classified correctly.

 *Mistakes on other examples may steer the parameters in different directions, so will the algorithm **converge** to something useful if we repeatedly cycle through the training examples?*

The Perceptron Algorithm - Convergence


 Yes, it will converge in the **realizable** case.

Theorem 2.1 *The perceptron update rule converges after a finite number of mistakes when the training examples are linearly separable through origin.* ← Perceptron convergence theorem

Note:

- The finite number of steps could still be **many**
- There can be **more than one solution**
 - sensitive to initialization
- If the dataset is **not linearly separable**, the training will **never converge**

Exercise: Perceptron Algorithm



$$x^{(1)} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}, y^{(1)} = 1$$

$$x^{(2)} = \begin{bmatrix} -1 \\ 2 \end{bmatrix}, y^{(2)} = -1$$

$$x^{(3)} = \begin{bmatrix} 0 \\ -1 \end{bmatrix}, y^{(3)} = -1$$

→ Random weight initialization

$$\theta^{(0)} = \begin{bmatrix} 1 \\ -0.8 \end{bmatrix}$$

→ Check if the point is classified correctly:



$$\begin{aligned} y^{(1)}(\theta^{(0)} \cdot x^{(1)}) &= (+1)((1 * 1) + (-0.8 * 2)) \\ &= (1 - 1.6) = -0.6 < 0 \end{aligned}$$

→ Update $\theta^{(0)}$:


$$\theta^{(1)} = \theta^{(0)} + y^{(1)}x^{(1)} = \begin{bmatrix} 1 \\ -0.8 \end{bmatrix} + \begin{bmatrix} 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 2 \\ 1.2 \end{bmatrix}$$

Perceptron update rule:

$$\text{if } y^{(t)} \neq h(x^{(t)}; \theta^{(k)}) \text{ then}$$
$$\theta^{(k+1)} = \theta^{(k)} + y^{(t)}x^{(t)}$$



Exercise: Perceptron Algorithm



Input data


$$x^{(1)} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}, y^{(1)} = 1$$

$$x^{(2)} = \begin{bmatrix} -1 \\ 2 \end{bmatrix}, y^{(2)} = -1$$

$$x^{(3)} = \begin{bmatrix} 0 \\ -1 \end{bmatrix}, y^{(3)} = -1$$


→ **1st Update**

$$\theta^{(1)} = \begin{bmatrix} 2 \\ 1.2 \end{bmatrix}$$




→ **2nd Update**

$$\theta^{(2)} = \begin{bmatrix} 3 \\ -0.8 \end{bmatrix}$$




→ **3rd Update**

$$\theta^{(3)} = \begin{bmatrix} 3 \\ 0.2 \end{bmatrix}$$




Exercise: Perceptron Algorithm

- Perceptron algorithm **oscillates** and **terminates with zero error** in **linearly separable** case.



Exercise: Perceptron Algorithm




Input data

$$x^{(1)} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}, y^{(1)} = 1$$

$$x^{(2)} = \begin{bmatrix} -1 \\ 2 \end{bmatrix}, y^{(2)} = -1$$

$$x^{(3)} = \begin{bmatrix} 0 \\ -1 \end{bmatrix}, y^{(3)} = -1$$

→ **Random weight initialization** →
(different initialized weights)



different decision boundaries


Linear Classifier (with offset)

Linear Classifier with offset

- We extend the set of linear classifiers slightly by **including a scalar offset parameter θ_0** .
 - This parameter will **enable us to place the decision boundary anywhere in \mathbb{R}^d** , not only through the origin.
- **Linear classifier with offset** is defined as:

$$h(x; \theta, \theta_0) = \text{sign}(\theta \cdot x + \theta_0) = \begin{cases} +1, & \theta \cdot x + \theta_0 \geq 0 \\ -1, & \theta \cdot x + \theta_0 < 0 \end{cases}$$

- The hyper-plane (line in 2d) $\theta \cdot x + \theta_0 = 0$ is **oriented parallel** to $\theta \cdot x = 0$.
- Vector θ is still **orthogonal** to the decision boundary.
- And θ also still defines the **positive direction** in the sense that if we move x in this direction, the value of $\theta \cdot x + \theta_0$ increases.



Linear Classifier with offset

- Both linear separability and the perceptron algorithm for learning linear classifiers **generalize easily** to the case of linear classifiers with offset.
- Linearly Separable**

Definition 3.1 Training examples $S_n = \{(x^{(t)}, y^{(t)}), t = 1; \dots, n\}$ are *linearly separable* if there exists a parameter vector $\hat{\theta}$ and offset parameter $\hat{\theta}_0$ such that $y^{(t)}(\hat{\theta} \cdot x^{(t)} + \hat{\theta}_0) > 0$ for all $t = 1, \dots, n$.

- Perceptron algorithm**

- Random initialize $\theta^{(0)}$ (vector) and $\theta_0^{(0)}$ (scalar)
- Perceptron update rule: if $y^{(t)} \neq h(x^{(t)}; \theta^k, \theta_0^k)$ then

$$\theta^{(k+1)} = \theta^{(k)} + y^{(t)} x^{(t)}$$


$$\theta_0^{(k+1)} = \theta_0^{(k)} + y^{(t)}$$

Perceptron convergence theorem


If training examples are linearly separable (**not necessarily through the origin**), then the perceptron algorithm converges after a finite number of mistakes.

Exercise: Linear Classifier with offset


- If training examples are **linearly separable through origin**, they are **clearly linearly separable (see last definition)**. Is the converse true?






(a)



(b)



(c)



Summary

1. What is a linearly separable set of examples?

Definition 3.1 Training examples $S_n = \{(x^{(t)}, y^{(t)}), t = 1; \dots, n\}$ are linearly separable if there exists a parameter vector $\hat{\theta}$ and offset parameter $\hat{\theta}_0$ such that $y^{(t)}(\hat{\theta} \cdot x^{(t)} + \hat{\theta}_0) > 0$ for all $t = 1, \dots, n$.

Definition 90 (Linear separability). If all the training data for class 1 lies on one side of a hyperplane, and for class 0 on the other, the data is said to be linearly separable.

From "Bayesian Reasoning and Machine Learning"

2. How the Perceptron algorithm works?

Perceptron update rule (through origin):

if $y^{(t)} \neq h(x^{(t)}; \theta^{(k)})$ then

$$\theta^{(k+1)} = \theta^{(k)} + y^{(t)}x^{(t)}$$

Perceptron update rule (general case):

if $y^{(t)} \neq h(x^{(t)}; \theta^k, \theta_0^k)$ then

$$\theta^{(k+1)} = \theta^{(k)} + y^{(t)}x^{(t)}$$

$$\theta_0^{(k+1)} = \theta_0^{(k)} + y^{(t)}$$

3. What is the guarantee of the Perceptron algorithm when the dataset is linearly separable?

Converges after a finite number of mistakes!

← Perceptron convergence theorem

Acknowledgements

- A lot of slides and content of this lecture are adopted from:
 - McGill COMP-652 Machine Learning
 - SUTD 50.007 Machine Learning, Spring 2023 (Asst Prof. Malika Meghjani)
 - MIT 6.036 Introduction to Machine Learning