


50.007 Machine Learning

2026 Spring

3. Hinge Loss

Na Zhao
Assistant Professor, ISTD

Recap



Recap - Linear Classifier (separable case)

1. Training Set ([Linearly Separable](#))

$$(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(n)}, y^{(n)})$$

2. Model ([Linear Classifier](#))

$$h(x; \theta) = \text{sign}(\theta_1 x_1 + \dots + \theta_d x_d)$$

3. Training Error ([Zero-one Loss](#))

$$\epsilon_n(\theta) = \frac{1}{n} \sum_{(x,y) \in \mathcal{S}_n} \mathbb{I}[y(\theta^\top x) \leq 0]$$

4. Algorithm ([The Perceptron Algorithm](#))

if $y^{(t)} \neq h(x^{(t)}; \theta^{(k)})$ then

$$\theta^{(k+1)} = \theta^{(k)} + y^{(t)} x^{(t)}$$

Recap - Linear Classifier (separable case)

- **Zero-one loss:**

$$\mathcal{E}_n(\theta) = \frac{1}{n} \sum_{t=1}^n [\![y^{(t)} \neq h(x^{(t)}; \theta)]\!] = \frac{1}{n} \sum_{t=1}^n [\![y^{(t)}(\theta \cdot x^{(t)}) \leq 0]\!]$$

- **Realizable case:** *zero* training error

- **Linear classifier with offset:**

$$h(x; \theta, \theta_0) = \text{sign}(\theta \cdot x + \theta_0) = \begin{cases} +1, & \theta \cdot x + \theta_0 \geq 0 \\ -1, & \theta \cdot x + \theta_0 < 0 \end{cases}$$

- **Perceptron Update Rule:**


if $y^{(t)} \neq h(x^{(t)}; \theta^k, \theta_0^k)$ then

$$\theta^{(k+1)} = \theta^{(k)} + y^{(t)}x^{(t)}$$

$$\theta_0^{(k+1)} = \theta_0^{(k)} + y^{(t)}$$

- + **Perceptron convergence theorem**

If training examples are **linearly separable**, then the perceptron algorithm **converges** after a finite number of mistakes.



θ is **orthogonal** to the decision boundary.
 θ **points in direction** of region labelled **+1**.

Learning Objectives

You should be able to know:

1. What is a convex function and how to check if a function is convex or not?
2. What is hinge loss?
3. How is hinge loss different from zero-one loss?
4. What is stochastic gradient descent?
5. How to implement the stochastic (sub-)gradient descent?


global min pt - lowest possible point anywhere
local min pt- lowest pon nearby

Convex Function

A function $f(x)$ is **convex** if for any two points x_1 and x_2 in the domain of $f(x)$, and for any t in the range **[0,1]**, the following condition must holds true:

$$f(tx_1 + (1 - t)x_2) \leq tf(x_1) + (1 - t)f(x_2)$$

- This means: line segment between any pair of points on the curve of f to be **above or just meets** the graph.
- Intuitively: the function is **bowl-shaped**.
- Convex function ensures that there is **only one global minimum**, and **no local minimum**, make optimization easier and more reliable.




Graph of a convex function : https://en.wikipedia.org/wiki/Convex_function


Non-convex Function

- A function is said to be **non-convex** if it is not convex.
- Geometrically, a non-convex function is one that **curves downwards** or has **multiple peaks and valleys**.
- The **challenge** with optimizing non-convex functions: may **get stuck in a local minimum (or saddle point)** and miss the global minimum, which is the optimal solution we are looking for.

saddle point - gradient = 0, not min or max




Graph of a non-convex function ([source](#))



Graph of a non-convex function with Local minima, Global minima, and a Saddle point ([source](#))


Convex Function Examples



A non-negative sum of convex functions is convex

Linear Classifier – Non-Separable Case

- If training examples are **not linearly separable**, the **perceptron algorithm** will **not converge** nor will it find the classifier with the **smallest error**.
- **Challenge:** The 0-1 risk/loss is a **non-convex** objective, and hence is **hard to optimize** (NP-hard) in this **non-realizable case**.



$$\text{Loss } h = \ell_h(z) = \begin{cases} 1 - y(\theta^\top x) & z \leq 0 \\ 0 & z > 0 \end{cases}$$

Hinge Loss

- As convex problems can be learned efficiently, we can **upper bound** the non-convex loss function by a **convex surrogate loss** function.
- Hinge loss** is a convex surrogate for the zero-one loss!

- **Zero-one loss**

$$\text{Loss}_{0|1}(z) = \mathbb{I}[y(\theta^\top x) \leq 0]$$


- **Hinge loss**

$$\text{Loss}_h(z) = \max\{1 - y(\theta^\top x), 0\}$$

CONVEX!

Penalize larger mistakes more.

Penalize near-mistakes, i.e. $0 \leq y(\theta^\top x) \leq 1$.



Zero-one Loss ~~VS~~ Hinge Loss - Examples

■ Zero-one loss

$$\text{Loss}_{0|1}(z) = \mathbb{I}[y(\theta^T x) \leq 0]$$

Example 1

$$z = y(\theta^T x) = -1(0.4) = -0.4 \Rightarrow \text{means zero-loss}$$

- original label = -1 and prediction score = 0.4 (*this means the model predicted class as 1*)
- penalty = $\max(0, 1+1(0.4)) = 1.4$ which is higher than penalty=1 if we use 0-1 loss.

Example 2 \nearrow same label $z = y(\theta^T x) = 1(-0.9) = -0.9 < 0$, indicator return 1

- original label = 1 and prediction score = (-0.9) (*this means the model predicted class as -1*)
- penalty = $\max(0, 1-1(-0.9)) = 1.9$ which is a very high penalty since the prediction was inaccurate
↳ hinge loss

Example 3

- original label = 1 and prediction score = 0.7 (*this means the model predicted class as 1*)
- penalty = $\max(0, 1-1(0.7)) = 0.3$ where **loss is low but not 0**, since the prediction is correct yet the prediction score is not very high (less than 1)

$$\begin{aligned} \text{eg 4: } y &= 1, \text{ pred score} = 1.4 \\ z &= 1(1.4) = 1.4 > 0 \end{aligned}$$

Empirical Risk Minimization

- **Empirical risk** is the training error of all training samples:

$$R_n(\theta) = \frac{1}{n} \sum_{t=1}^n \text{Loss}(y^{(t)}(\theta \cdot x^{(t)}))$$

can be any loss function (hinge / zero-one)

- If we adopt **hinge loss**, it becomes:


$$R_n(\theta) = \frac{1}{n} \sum_{t=1}^n \text{Loss}_h(y^{(t)}(\theta \cdot x^{(t)})) = \frac{1}{n} \sum_{t=1}^n \max\{1 - y^{(t)}(\theta \cdot x^{(t)}), 0\}$$

- Remember that the classifier training **goal** is to **minimize** the test or **generalization error**.
- The **idea** of using hinge loss is to give the classifier a little more **feedback** in terms of how **close** its predictions are to the **training labels**, thus it will **generalize** better.
- Moreover, now $R_n(\theta)$ is a **convex** function, and convexity of empirical risk allows us to find the **minimum** even in **non-realizable** case (the examples are *not linearly separable*).


If you want to learn more about (non)convex optimization:

- Theory: https://web.stanford.edu/~boyd/cvxbook/bv_cvxbook.pdf
- Funny examples: <http://www.stat.cmu.edu/~ryantibs/convexopt/lectures/nonconvex.pdf>

Where is the Godzilla?



Where are the Godzillas?



Each Godzilla defines a local minima and the “heaviest”
Godzilla: The global minima “non-convex”

Gradient Descent

- Gradient descent is a first-order **iterative** algorithm for finding a **local minimum** (global minimum if the function is convex) of a **differentiable** multivariate function.
- When use **gradient descent** to minimize $R_n(\theta)$, it simply iteratively updates the parameters:

$$\theta^{(k+1)} = \theta^{(k)} - \eta_k \nabla_{\theta} R_n(\theta)_{\theta=\theta^{(k)}}$$

- η_k is known as the *step-size* or *learning rate*.
- The gradient points in the **direction** where $R_n(\theta)$ **increases**.

$$\nabla_{\theta} R_n(\theta) = \left[\frac{\partial R_n(\theta)}{\partial \theta_1}, \dots, \frac{\partial R_n(\theta)}{\partial \theta_d} \right]^T$$

- Thus, the weight is updated in the **opposite direction** to minimize error.




Illustration of gradient descent on a series of level sets

Stochastic Gradient Descent

- In **Gradient Descent**, the **entire training dataset** is used to compute the **gradient**.
- **Stochastic Gradient Descent (SGD)** is a stochastic approximation of gradient descent, as it replaces the **actual gradient** by an **estimate** calculated from a **randomly selected sample**.
 - Reduce computational cost *→ update may be in wrong direction*
 - Converge faster (but with more erratic updates)
- The procedure of SGD algorithm:

```
 $\theta^{(0)} = 0$  (vector)      # Random initialization  
select  $t \in \{1, \dots, n\}$  at random  
 $\theta^{(k+1)} = \theta^{(k)} - \eta_k \nabla_{\theta} \text{Loss}_h(y^{(t)} \theta \cdot x^{(t)})|_{\theta=\theta^{(k)}}$  # Update  
Repeat the update step until stopping criterion is met.
```

Are we done here?

NO!!!

■ Hinge loss

$$\text{Loss}_h(z) = \max\{1 - y(\theta^T x), 0\}$$


- The hinge loss function is defined as **piecewise linear function**, which is **not everywhere differentiable!** 😞

Sub-differential Problem


- $R_n(\theta)$ is also **not differentiable everywhere** as it is a **summation** of hinge loss functions.

$$R_n(\theta) = \frac{1}{n} \sum_{t=1}^n \text{Loss}_h(y^{(t)}(\theta \cdot x^{(t)})) = \frac{1}{n} \sum_{t=1}^n \max\{1 - y^{(t)}(\theta \cdot x^{(t)}), 0\}$$

- There are **several possible gradients** at the **kinks** (*points where the function is not differentiable*) which are collectively defined as **sub-differential**.
- In our case, to minimize $R_n(\theta)$, we need to select **one possible gradient** at **each kink** *regardless of how many choices there are*.



- ❖ Recall and observe the piece-wise linear shape of hinge loss function



[A sample representation of sub-differential](#)

Stochastic Sub-Gradient Descent


- Stochastic Gradient Descent

$$\theta^{(0)} = 0 \text{ (vector)} \quad \# \text{Random initialization}$$

select $t \in \{1, \dots, n\}$ at random,

$$\theta^{(k+1)} = \theta^{(k)} - \eta_k \nabla_{\theta} \text{Loss}_h(y^{(t)} \theta \cdot x^{(t)})|_{\theta=\theta^{(k)}} \quad \# \text{Update}$$

Repeat the update step until stopping criterion is met.



If the agreement $y^{(t)}(\theta^{(k)} x^{(t)}) > 1$ then the **loss** is identically zero and so is the **gradient**. \Rightarrow **No update**

Else: $\nabla_{\theta} \text{Loss}_h(y^{(t)} \theta \cdot x^{(t)})|_{\theta=\theta^{(k)}} = \nabla_{\theta}(1 - y^{(t)} \theta \cdot x^{(t)})|_{\theta=\theta^{(k)}} = -y^{(t)} x^{(t)}$

y *gradient*

- Stochastic Sub-Gradient Descent

$$\theta^{(0)} = 0 \text{ (vector)} \quad \# \text{Random initialization}$$

select $t \in \{1, \dots, n\}$ at random,

$$\text{if } y^{(t)}(\theta^{(k)} \cdot x^{(t)}) \leq 1, \text{ then } \theta^{(k+1)} = \theta^{(k)} + \eta_k y^{(t)} x^{(t)} \quad \# \text{Update}$$

update only if this applicable

Repeat the update step until stopping criterion is met.

perceptron algorithm
Does it look familiar?

Stochastic Sub-Gradient Descent

$\theta^{(0)} = 0$ (vector) # Random initialization ①

select $t \in \{1, \dots, n\}$ at random, # Random select training data from sample
 if $y^{(t)}(\theta^{(k)} \cdot x^{(t)}) \leq 1$, then
 $\theta^{(k+1)} = \theta^{(k)} + \eta_k y^{(t)} x^{(t)}$ # Update

Repeat the update step until stopping criterion is met.

$$\max[1 - y(\theta \cdot x), 0]$$

- Differences
 - SSGD also penalizes (update parameters) **near** mistakes, while the perceptron **only** penalizes **mistakes**.
 - SSGD use a **decreasing** learning rate η_k (later updates will be smaller), while the perceptron sets it to **one**.
 - SSGD **randomly selects** the training example, while the perceptron **cycles** through all examples **in order**.

VS Perceptron Algorithm

```

① Random initialize  $\theta = 0$ 
while TRUE do
     $m = 0$ 
    for  $(x^{(t)}, y^{(t)}) \in D$  do ②
        if  $y^{(t)}(\theta \cdot x^{(t)}) \leq 0$  then
             $\theta \leftarrow \theta + y^{(t)} x^{(t)}$ 
             $m \leftarrow m + 1$ 
        end if
    end for
    if  $m = 0$  then
        break
    end if
end while

```

$$\max[-y(\theta \cdot x), 0]$$

How to select learning rate?

- Any **positive learning rate** that satisfies the following condition would permit the algorithm to **converge** though the speed may vary.

$$\sum_{k=1}^{\infty} \eta_k^2 < \infty, \quad \sum_{k=1}^{\infty} \eta_k = \infty$$

- In our context, $\eta_k = 1/(k + 1)$ would ensure that our stochastic sub-gradient descent algorithm converges to the minimum of $R_n(\theta)$.
- In practice, one can choose other learning rates based on **empirical observations**.
 - Simply setting $\eta_k = 0.1$ turns out to be a **popular choice** in many problems.

How to get the best parameters?

- Due to the **erratic updates** in SGD, the $\theta^{(k)}$ may **not be the best** solution obtained so far (till k -th step).
- Hence, it is often quite beneficial to **keep track** of the best solution $\theta^{(i_k)}$, where

$$i_k = \operatorname{argmin}_{i=1,\dots,k} R_n(\theta^{(i)})$$

- When increasing k , the empirical risk of $\theta^{(k)}$, i.e., $R_n(\theta^{(k)})$ goes down in a noisy fashion or oscillates, while $R_n(\theta^{(i_k)})$ is **monotonically decreasing** (*guaranteed by the definition!*)
- Note: $\lim_{k \rightarrow \infty} R_n(\theta^{(i_k)})$ is **not necessarily zero** as the points may be **not linearly separable**.

```
 $\theta^{(0)} = 0$  (vector)      # Random initialization  
select  $t \in \{1, \dots, n\}$  at random,  
if  $y^{(t)}(\theta^{(k)} \cdot x^{(t)}) \leq 1$ , then  
 $\theta^{(k+1)} = \theta^{(k)} + \eta_k y^{(t)} x^{(t)}$     # Update  
Repeat the update step until stopping criterion is met.
```

Linear Classifier (Non-separable case)

1. Training Set ([Not Necessarily Linearly Separable](#))

$$(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(n)}, y^{(n)})$$

2. Model ([Linear Classifier](#))

$$h(x; \theta) = \text{sign}(\theta_1 x_1 + \dots + \theta_d x_d)$$

3. Training Error ([Hinge Loss](#))

$$R_n(\theta) = \frac{1}{n} \sum_{(x,y) \in S_n} \max\{1 - y(\theta^\top x), 0\}$$

4. Algorithm ([Stochastic Sub-Gradient Descent](#))

select $t \in \{1, \dots, n\}$ at random,

if $y^{(t)}(\theta^{(k)} \cdot x^{(t)}) \leq 1$, then

$$\theta^{(k+1)} = \theta^{(k)} + \eta_k y^{(t)} x^{(t)}$$

update θ

Summary

1. What is a convex function and how to check if a function is convex or not?

$$f(tx_1 + (1-t)x_2) \leq tf(x_1) + (1-t)f(x_2)$$

Convex: Line segment between any pair of points on the curve of f to be above or just meets the graph.

2. What is hinge loss?

Hinge loss is a convex surrogate for the zero-one loss, defined as $\text{Loss}_h(z) = \max\{1 - y(\theta^\top x), 0\}$

3. How is hinge loss different from zero-one loss?

Hinge loss function is convex, it penalizes near mistake and penalizes larger mistakes more than 0-1 loss.

4. What is stochastic gradient descent?

SGD is a stochastic approximation of gradient descent by estimating the gradient from a randomly selected sample.

5. How to implement the stochastic (sub-)gradient descent?

If the agreement $y^{(t)}(\theta^{(k)}x^{(t)}) > 1$ then the **loss** is **identically zero** and so is the **gradient**. \Rightarrow **No update**

Else: $\nabla_{\theta} \text{Loss}_h(y^{(t)}\theta \cdot x^{(t)})|_{\theta=\theta^{(k)}} = \nabla_{\theta}(1 - y^{(t)}\theta \cdot x^{(t)})|_{\theta=\theta^{(k)}} = -y^{(t)}x^{(t)}$

Acknowledgements

- Some slides and content of this lecture are adopted from:
 - McGill COMP-652 Machine Learning
 - SUTD 50.007 Machine Learning, Spring 2023 (Asst Prof. Malika Meghjani)
 - MIT 6.036 Introduction to Machine Learning
 - Shalev-Shwartz, Shai, and Shai Ben-David. "Understanding machine learning: From theory to algorithms". Cambridge university press, 2014. (Chapter 8.2 & 12.3)